Machine Learning 2 - Cheat Sheet
Andrea Jemmett - April 17, 2016

1 Probability Theory

1.1 Independence

$p(X,Y) = p(X)p(Y) \Leftrightarrow p(X|Y) = p(X) \Leftrightarrow p(Y|X) = p(Y)$

1.2 Conditional Independence

$X \perp\!\!\!\perp Y \mid Z \Longleftrightarrow p(X,Y|Z) = p(X|Z)p(Y|Z)$

1.3 Sum and Product Rules

$p(X,Y) = p(X)p(Y|X), \quad p(X,Y,Z) = p(X)p(Y|X)p(Z|X,Y)$
$p(X) = \sum_Y p(X,Y)$

1.4 Bayes' Theorem

$p(Y|X) = \frac{p(X|Y)p(Y)}{p(X)}, \quad p(Y|X,Z) = \frac{p(X|Y,Z)p(Y|Z)}{p(X|Z)}$

2 Distributions

Binary	Bernoulli	Binomial	Beta
Discrete	Categorical	Multinomial	Dirichlet

2.1 Bernoulli Distribution

$\mathrm{Ber}(x|n) = \mu^x(1-\mu)^{1-x}, \quad \mathbb{E}[x] = \mu, \quad Var[x] = \mu - \mu^2,$
$P(D,\mu) = \prod_{n=1}^N \mu^{x_n}(1-\mu)^{1-x_n}, \quad \mu_{\mathrm{ML}} = \frac{1}{N}\sum_{n=1}^N x_n$

2.2 Binomial Distribution

$\mathrm{Bin}(m|N,\mu) = \binom{N}{m}\mu^m(1-\mu)^{N-m}, \quad \frac{n!}{k!(n-k)!} = \binom{n}{k},$
$\mathbb{E}[m] = N\mu, \quad Var[m] = N\mu(1-\mu), \quad \mu_{\mathrm{ML}} = \frac{m}{N}$

2.3 Categorical Distribution

$p(\mathbf{x}|\boldsymbol{\mu}) = \prod_k \mu_k^{x_k}, \quad \boldsymbol{\mu} \in \{0,1\}^K, \quad \sum_k \mu_k = 1, \quad \boldsymbol{\mu}_{\mathrm{ML}} = \frac{\mathbf{m}}{N},$
$m_k = \sum_n x_{nk}, \quad \mathrm{Mult}(m_1 \ldots, m_k|N,\boldsymbol{\mu}) = (\frac{N!}{m_1!,\ldots,m_k}\prod_k)\mu_k^{mk}$

2.4 Beta Distribution

$\mathrm{Beta}(\mu|a,b) = \frac{\Gamma(a+b)}{\Gamma(a)\Gamma(b)}\mu^{a-1}(1-\mu)^{b-1},$ [1] $\quad \mathbb{E}[\mu] = \frac{a}{a+b},$
$Var[x] = \frac{ab}{(a+b)^2(a+b+1)}, \quad p(\mu|m,l,a,b) \propto \mu^{m+a-1}(1-\mu)^{l+b-1}$

2.5 Gamma Distribution

$\mathrm{Gamma}(\tau|a,b) = \frac{b^a}{\Gamma(a)}\tau^{a-1}e^{-b\tau}, \quad \mathbb{E}[\tau] = \frac{a}{b}, \quad Var[\tau] = \frac{a}{b^2},$
$\mathrm{mode}[\tau] = \frac{a-1}{b}$ for $a \geq 1$, $\quad \mathbb{E}[\ln\tau] = \psi(a) - \ln b,$
$H(\tau) = ln\Gamma(a) - (a-1)\psi(a) - \ln b + a$

2.6 Multinomial Distribution

$\mathbf{x} = [0,0,0,0,1,0,0]^\intercal, \quad \sum_{k=1}^K x_k = 1, \quad p(\mathbf{x}|\boldsymbol{\mu}) = \prod_{k=1}^K \mu_k^{x_k},$
$\sum_{k=1}^K \mu_k = 1, \quad \mu_k^{\mathrm{ML}} = \frac{m_k}{N}, \quad m_k = \sum_{k=1}^K x_{nk}$

2.7 Dirichlet Distribution

$\mathrm{Dir}(\boldsymbol{\mu}|\boldsymbol{\alpha}) = \frac{\Gamma(\alpha_0)}{\Gamma(\alpha_1)\cdots\Gamma(\alpha_k)}\prod_{k=1}^K \mu_k^{\alpha_k-1}, \quad \alpha_0 = \sum_{k=1}^K \alpha_k$

2.8 Gaussian Distribution

$\mathcal{N}(x|\mu,\sigma) = \frac{1}{\sqrt{2\pi\sigma^2}}\exp(-\frac{1}{2\sigma^2}(x-\mu)^2),$
$\mathcal{N}(\mathbf{x}|\boldsymbol{\mu},\boldsymbol{\Sigma}) = (2\pi)^{-\frac{D}{2}}|\boldsymbol{\Sigma}|^{-\frac{1}{2}}\exp\left\{-\frac{1}{2}(\mathbf{x}-\boldsymbol{\mu})^\intercal\boldsymbol{\Sigma}^{-1}(\mathbf{x}-\boldsymbol{\mu})\right\}$

2.8.1 ML for the Gaussian

$\ln p(X|\boldsymbol{\mu},\boldsymbol{\Sigma}) = -\frac{ND}{2}\ln(2\pi) - \frac{N}{2}\ln|\boldsymbol{\Sigma}| - \frac{1}{2}\sum_{n=1}^N(\mathbf{x}_n-\boldsymbol{\mu})^\intercal\boldsymbol{\Sigma}^{-1}(\mathbf{x}_n-\boldsymbol{\mu}),$
$\boldsymbol{\mu}_{\mathrm{ML}} = 1/N\sum_{n=1}^N \mathbf{x}_n, \quad \boldsymbol{\Sigma}_{\mathrm{ML}} = 1/N\sum_{n=1}^N(\mathbf{x}_n-\boldsymbol{\mu})^\intercal(\mathbf{x}_n-\boldsymbol{\mu})$

2.8.2 Stochastic gradient descent Gaussian

$\max\ P(x_1,\cdots,x_n|\theta),\ \theta^N = \theta^{N-1} + \alpha_{N-1}\frac{\partial}{\partial\theta^{N-1}}\ln\ p(x_n|\theta^{N-1})$

2.8.3 Marginal and Conditional Gaussians

Given $p(\mathbf{x}) = \mathcal{N}(\mathbf{x}|\boldsymbol{\mu},\boldsymbol{\Lambda}^{-1})$ and $p(\mathbf{y}|\mathbf{x}) = \mathcal{N}(\mathbf{y}|\mathbf{A}\mathbf{x}+\mathbf{b},\mathbf{L}^{-1})$. We get
$p(\mathbf{y}) = \mathcal{N}(\mathbf{y}|\mathbf{A}\boldsymbol{\mu}+\mathbf{b},\mathbf{L}^{-1}+\mathbf{A}\boldsymbol{\Lambda}^{-1}\mathbf{A}^\intercal)$ and
$p(\mathbf{x}|\mathbf{y}) = \mathcal{N}(\mathbf{x}|\boldsymbol{\Sigma}[\mathbf{A}^\intercal\mathbf{L}(\mathbf{y}-\mathbf{b})+\boldsymbol{\Lambda}\boldsymbol{\mu}],\boldsymbol{\Sigma})$
where $\boldsymbol{\Sigma} = (\boldsymbol{\Lambda}+\mathbf{A}^\intercal\mathbf{L}\mathbf{A})^{-1}$.

2.9 Student's T distribution

The heavy tail of the student-t distribution makes it more robust against outliers.
$St(x|\mu,\lambda,\nu) = \frac{\Gamma(\nu/2+1/2)}{\Gamma(\nu/2)}(\frac{\lambda^{1/2}}{(\pi\nu)^{D/2}})[1+\frac{\lambda(x-\mu)^2}{\nu}]^{-\nu/2-D/2},$
$f_x(x) = \frac{\Gamma[(\nu+p)/2]}{\Gamma(\nu/2)\nu^{p/2}\pi^{p/2}|\Sigma|^{1/2}[1+1/\nu(x-\mu)^T\Sigma^{-1}(x-\mu)]^{(\nu+p)/2}}$
$\mathbb{E}(\mathbf{x}) = \frac{\Gamma(D/2+v/2)}{\Gamma(v/2)}\frac{|\Lambda|^{1/2}}{(\pi v)^{D/2}}\int[1+\frac{(x-\mu)^T\Lambda(x-\mu)}{v}]^{-D/2-v/2}x\mathrm{d}x$

3 Graphical Models

Capture noise, for reasoning, enormous datasets, for causality, for designing models, CIR are encoded in the graph, for inference.

3.1 Directed GMs a.k.a. Bayesian Networks

Nodes are connected by *directed* arrows. The full joint distribution is $p(\mathbf{x}) = \prod_{k=1}^K p(x_k|\mathrm{pa}(x_k))$
Directed Acyclic Graphs are BNs without *directed* loops.

3.1.1 Blocking Rules

$\bigcirc \rightarrow \bullet \rightarrow \bigcirc \qquad \bigcirc \leftarrow \bullet \rightarrow \bigcirc \qquad \bigcirc \rightarrow \cdots \bigcirc \cdots \leftarrow \bigcirc$

3.1.2 D-separation

$A \perp\!\!\!\perp B \mid C$ holds if each path that connects a node in A with a node in B is *blocked*, that is
a) the arrows on the path meet either head-to-tail or tail-to-tail at the node and the node is in C, *or*
b) the arrows meed head-to-head at the node and either the node nor any of its descendants is in C.

3.2 Markov Random Fields

Graphical Models with undirected edges (a.k.a. Undirected Graphical Models).

3.2.1 Conditional Independence

$A \perp\!\!\!\perp B \mid C$ holds if all paths connecting every node in A to every other node in B is 'blocked' by a node in C (blocked means passing through).

3.2.2 Cliques and Maximal Cliques

A *clique* is a subset of nodes in the graph such that there exists a link between all pairs of nodes in the subset (the set of nodes in a clique is fully connected). A *maximal clique* is a clique such that it is not possible to include any other nodes from the graph in without it ceasing to be a clique.

3.2.3 Factorization

A MRF can be factorized using *potential functions* over its maximal cliques: $p(\mathbf{x}) = \frac{1}{Z}\prod_C \psi_C(\mathbf{x}_C) \quad Z = \sum_x \prod_C \psi_C(\mathbf{x}_C)$

3.2.4 Relation to Directed Graphs

To transform a directed graph into an undirected one we have to perform a *moralization* process of "marrying the parents" of each node (by linking them) and removing all remaining arrows.

3.3 Markov Blanket

$p(x_i|x_{\mathrm{MB}_i}, x_{\mathrm{rest}}) = p(x_i|x_{\mathrm{MB}_i})$
Directed case: The MB of x_i consists of: the parents of x_i, the children of x_i and the co-parents of the children of x_i.
Undirected case: All neighboring nodes of x_i.

3.4 Naive Bayes

The problem of finding a label y^* for some previously unobserved vector of features $\mathbf{x}^*$, while having previously observed $\{\mathbf{x}_n, y_n\}_{n=1,\ldots,N}$. We build a *generative* GM, linking the label y_n with each feature of the feature vector $\mathbf{x}_n$. This implies that all features are independent of each other given the label. For a single data case the joint probability is $p(y, x_1, \ldots, x_D|\boldsymbol{\eta},\boldsymbol{\theta}_i) = p(y|\boldsymbol{\eta})\prod_{i=1}^D p(x_i|\boldsymbol{\theta}_i)$
while the probability of the full dataset is $\prod_{n=1}^N p(y_n|\boldsymbol{\eta})\prod_{i=1}^D p(x_i|\boldsymbol{\theta}_i)$
Note that this generative form is chosen as the other option (features pointing towards the label) would imply a very highly parameterized model, as we would considering $p(y|x_1,\ldots,x_D)$.
Finding a label y^* implies finding
$y^* = \mathrm{argmax}_y\left\{\ln p(y|\boldsymbol{\eta}) + \sum_{i=1}^D \ln p(x_i^*|y,\boldsymbol{\theta}_i)\right\}$

3.5 Maximum Likelihood Training in BNs

It is fast because the log-likelihood *decomposes* into a sum over all variables X_i. Learning all parameters reduces into a collection of independent tasks of learning $p(x_i|\mathrm{pa}_{x_i})$.
$p(x_i|\mathrm{pa}_{x_i}) = \frac{N(x_i,\mathrm{pa}_{x_i})}{N(\mathrm{pa}_{x_i})}$
is the number of times x_i co-occurred with pa_{x_i} divided by the number of times pa_{x_i} occurred.

3.6 Inference in Graphical Models

In which some of the variables are observed and we wish to compute the posterior distribution of one or more subsets of other variables.

3.6.1 Inference on a chain

$p(\mathbf{x}) = p(x_1,\ldots,x_N)$
$= \frac{1}{Z}\psi_{1,2}(x_1,x_2)\ldots\psi_{N-1,N}(x_{N-1},x_N)$
$= \sum_{x_1}\sum_{x_2}\cdots\sum_{x_{N-1}}\sum_{x_N} p(\mathbf{x})$
$= \sum_{x_1}\cdots\sum_{x_{n-1}}\sum_{x_{n+1}}\cdots\sum_{x_{N-1}}\sum_{x_N} p(\mathbf{x})$
$= \frac{1}{Z}\sum_{x_1}\cdots\sum_{x_{n-1}}\psi_{x_1,x_2}\cdots\psi_{x_{n-1},x_n}\mu_\beta(x_n)$
$\mu_\beta(x_n) = \frac{1}{Z}\sum_{x_{n+1}}\psi_{x_n,x_{n+1}}\cdots\sum_{x_N}\psi_{x_{N-1},x_N}$
$\mu_\alpha(x_n) = \frac{1}{Z}\sum_{x_{n-1}}\psi_{x_{n-1},x_n}\cdots\sum_{x_1}\psi_{x_2,x_1}$
$p(x_n) = \frac{1}{z}\mu_\alpha(x_n)\mu_\beta(x_n) \qquad O(NK^2)$

3.6.2 Factor Graphs

A tree is a graph with no loops. Both directed and undirected trees can be converted to a factor graph tree, but a directed tree could result in a non-tree structure when converted to an undirected representation. It is called a poly-tree (and not simply a tree) since its undirected representation (middle graph) includes a loop. The factor graph representation is again a tree. Factor graphs are the most general representation, and since any other tree representation can be easily converted to a factor tree, the sum-product algorithm is defined for factor trees.

[1] $\Gamma(x) = \int_0^1 u^{x-1}e^{-u} = 1, \quad \Gamma(x+1) = \Gamma(x)x, \quad \Gamma(x+1) = x!$

3.6.3 Sum-product algorithm

Probability of the factor graph: $p(\vec{x}) = \frac{1}{Z} \prod_\alpha f_\alpha(\vec{x}_\alpha)$

factor → variable message $\mu_{\alpha \to i}(x_i) = \sum\limits_{x_{\alpha \setminus i}} f_\alpha(\vec{x}_\alpha) \prod\limits_{j \in \alpha \setminus i} \mu_{j \to \alpha}$

variable → factor message $\mu_{j \to \alpha}(x_j) = \prod\limits_{\beta \in \mathrm{ne}(j) \setminus \alpha} \mu_{\beta \to j}(x_j)$

leaf node messages	x_l	is a leaf node:	$\mu_{l \to \delta}(x_l) = 1$
	ε	is a leaf node:	$\mu_{\varepsilon \to k}(x_k) = f_\varepsilon(x_k)$